

## A Supplementary Material – Correctness

In this section, we prove Oak's correctness. Since the rebalancer is orthogonal to our contribution, we omit it from the discussion of Oak's correctness. We only assume that RB1-3 hold. We note that a similar rebalance was fully proven in [17].

### A.1 Preliminaries

We consider a shared memory system consisting of a collection of shared variables accessed by threads, which also have local variables. An *algorithm* defines the behaviors of threads as deterministic state machines, where state transitions are associated with either an instance of a shared variable primitive (read, write, CAS, etc.) or a local step affecting the thread's local variables. A *configuration* describes the current state of all local and shared variables. An *initial configuration* is one where all variables hold an initial value. A data structure implementation provides a set of operations, each with possible parameters. We say that operations are *invoked* and *return* or *respond*. The invocation of an operation leads to the execution of an algorithm by a thread. Both the invocation and the return are local steps of a thread. A *run* of algorithm  $\mathcal{A}$  is an alternating sequence of configurations and steps, beginning with some initial configuration, such that configuration transitions occur according to  $\mathcal{A}$ . We say that two operations are *concurrent* in a run  $r$  if both are invoked in  $r$  before either returns. We use the notion of time  $t$  during a run  $r$  to refer to the configuration reached after the  $t^{\text{th}}$  step in  $r$ . An *interval* of a run  $r$  is a sub-sequence that starts with a step and ends with a configuration. The *interval of an operation*  $op$  starts with the invocation step of  $op$  and ends with the configuration following the return from  $op$  or the end of  $r$ , if there is no such return.

An implementation of concurrent data structure is *linearizable* [33] (a correctness condition for concurrent objects) if it provides the illusion that each invoked operation takes effect instantaneously at some point, called the *linearization point* (l.p.), inside its interval. A *linearization* of a run  $r$  ( $lin(r)$ ) is the sequential run constructed by serially executing each operation at its l.p.

### A.2 Linearizability proof

**Definition 1.** If there is an entry  $e$  in Oak that points to key  $k$  and handle  $h$ , (i.e.,  $\text{lookup}(k)$  returns  $e$  s.t.  $h = \text{handles}[\text{entries}[e].hi]$ ) and  $h.\text{deleted} = \text{false}$ , we say that  $h$  is *associated with*  $k$ .

**Claim 2.** If an Oak operation searches for key  $k$  and finds a non-deleted handle  $h$  ( $h.\text{deleted} = \text{false}$ ), then  $h$  is associated with  $k$ .

*Proof.* If an operation searches for  $k$  and finds  $h$ , then there is an entry  $e$  that points to  $k$ , since Oak ensures that there is at most one entry that points to  $k$ , and  $k$  is found only if there is such entry. This also means that  $e$  points to handle  $h$

(by handle index  $hi$ ). Assume that  $e$  does not point to handle  $h$ , then the handle index is now  $hi' \neq hi$ . If  $hi' = \perp$  then the handle index can be set only by a non-insertion operation using a CAS. According to Algorithm 3 this is only possible when  $h$  in  $\text{handles}[hi]$  is already deleted, but  $h$  is not deleted. Otherwise,  $hi' \neq hi$  and  $hi' \neq \perp$ , then the handle index can be set only by an insertion operation using a CAS. According to Algorithm 2 this is only possible when  $h$  in  $\text{handles}[hi]$  is already deleted, which is not the case. Therefore, there is an entry  $e$  that points to  $k$  and  $h$  and  $h.\text{deleted} = \text{false}$ , so by Definition 1  $h$  is associated with  $k$ .  $\square$

**Claim 3.** Assume handle  $h$  is associated with key  $k$  at time  $t$  in a run  $r$ . Then,  $h$  is associated with  $k$  at time  $(t + 1)$  in  $r$  if and only if the  $(t + 1)^{\text{st}}$  step in  $r$  is not the l.p. of a successful  $\text{remove}(k)$  operation.

*Proof.* Assume that  $h$  is not associated with  $k$  at time  $(t + 1)$ .

If there is no handle associated with  $k$  at time  $t + 1$ , then by Definition 1 either  $h.\text{deleted} = \text{true}$  or the entry's handle index ( $hi$ ) is  $\perp$ . In the first case, the only possible step that marks a handle as deleted is the l.p. of a successful  $\text{remove}(k)$ . In the second case, only non-insertion operations turn  $hi$  to  $\perp$  by using CAS (lines 56 and 74), and according to Algorithm 3 this is only possible when the handle is deleted. However, at time  $t$ ,  $h$  is still associated with  $k$ . Therefore, the entry's handle index ( $hi$ ) is not  $\perp$ .

Otherwise, there is a different handle  $h' \neq h$  that is associated with  $k$  at time  $t + 1$  ( $h' \neq \perp$ ). This change can only be done by an insertion operation using CAS (line 38). According to Algorithm 2 an insertion operation reaches that CAS only if the handle ( $h$ ) is already deleted (line 21). However, at time  $t$ ,  $h$  is still associated with  $k$ , and so there is no different handle that is associated with  $k$ .

Therefore, as long as the  $(t + 1)^{\text{st}}$  step is not the l.p. of a successful  $\text{remove}(k)$ , then  $h$  is still associated with  $k$  at time  $t + 1$  in  $r$ , and there is no handle associated with  $k$  at time  $t + 1$  if the  $(t + 1)^{\text{st}}$  step is a l.p. of a successful  $\text{remove}(k)$ , as required.  $\square$

**Claim 4.** Assume no handle is associated with key  $k$  at time  $t$  in a run  $r$ . Then, no handle is associated with  $k$  at time  $t + 1$  in  $r$  if and only if the  $(t + 1)^{\text{st}}$  step in  $r$  is not the l.p. of a successful insertion operation of  $k$ .

*Proof.* If no handle is associated at time  $t$ , and at time  $t + 1$  there is an associated handle, then according to Definition 1 either a handle's deleted flag turned from false to true, or the entry's handle index turned from  $\perp$  to a valid one. The former is not possible because the handles are initialized as not deleted and only become deleted by a remove; no operation turns a deleted handle to a non-deleted one. In the second case, this can only be done by a successful insertion operation, at its l.p. (line 38), as required.  $\square$

Look at the linearization  $lin(r)$  of run  $r$  using l.p.s defined in Section 4.5. From Claims 3 and 4, by induction on the steps of a run, we get:

**Corollary 5.** At any point in a concurrent run  $r$ , the set of keys associated with handles is exactly the same as the set of inserted keys and not removed keys, associated with the same handles, in  $lin(r)$  up to that point.

**Claim 6 (Get).** In run  $r$ , if  $get(k)$  returns  $h$  then the corresponding  $get(k)$  in  $lin(r)$  returns  $h$ , and if  $get(k)$  returns null then the corresponding  $get(k)$  in  $lin(r)$  returns null.

*Proof.* There are three cases for  $get$ 's l.p.:

1.  $get(k)$  finds a non-deleted handle  $h$  (line 6), then  $get(k)$  returns  $h$  and by Claim 2  $h$  is associated with  $k$ . By Corollary 5, in  $lin(r)$   $k$  is inserted and not removed (the map holds  $k$ ) and since this is the l.p. of  $get$  then the corresponding  $get(k)$  in  $lin(r)$  returns  $h$  as well.
2.  $Lookup(k)$  by  $get(k)$  (line 3) returns  $\perp$  or if  $get(k)$  reads that the handle index is  $\perp$  (line 4), then there is no handle associated with key  $k$ , and  $get(k)$  returns null. By Corollary 5, in  $lin(r)$  the map does not hold  $k$ , and since this is the l.p. of  $get$  then the corresponding  $get(k)$  in  $lin(r)$  returns null as well.
3.  $get(k)$  finds a deleted handle  $h$  at time  $t_2$  (line 6) and returns null. Then its l.p. is the later between the read of handle index  $hi$  by  $get(k)$  at time  $t_1 < t_2$  (line 4) and immediately after the set of deleted = true by  $remove(k)$  at some time  $t < t_2$ . Again there are two cases:
  - a. If  $t > t_1$  then the l.p. is immediately after the set of deleted = true then there is no handle associated with key  $k$ , and by Corollary 5, in  $lin(r)$  the map does not hold  $k$ , and the corresponding  $get(k)$  in  $lin(r)$  returns null as well.
  - b. If  $t_1 > t$  then the l.p. is the read of handle index  $hi$  by  $get(k)$  (line 4) at time  $t_1$ , after the set of deleted = true at time  $t$ . We need to show that at no time between  $t$  and  $t_1$  the handle index changed to  $hi' \neq hi$  and now it does not point to a deleted handle. Notice that only an insertion operation l.p. can change  $hi$  to  $hi'$ . Assume by contradiction that the l.p. of such an operation occurs between  $t$  and  $t_1$ . Then when  $get$  sees  $hi$  at time  $t_1$ , it is already  $hi'$  and not  $hi$ . A contradiction. Hence, at the l.p. of  $get(k)$ , there is no handle associated with key  $k$ , and by Corollary 5, in  $lin(r)$  the map does not hold  $k$ , so the corresponding  $get(k)$  in  $lin(r)$  returns null as required.  $\square$

**Claim 7 (PutIfAbsent).** In run  $r$ , if  $putIfAbsent(k)$  returns true then the corresponding  $putIfAbsent(k)$  in  $lin(r)$  returns true, and if  $putIfAbsent(k)$  returns false then in  $lin(r)$  the corresponding  $putIfAbsent(k)$  returns false.

*Proof.* If  $putIfAbsent(k)$  finds a non-deleted handle  $h$  (line 21), then  $putIfAbsent(k)$  returns false and by Claim 2  $h$  is associated with  $k$ . By Corollary 5, in  $lin(r)$   $k$  is inserted and not removed (the map holds  $k$ ) and since this is the l.p. of  $putIfAbsent$  then the corresponding  $putIfAbsent(k)$  in  $lin(r)$  returns false as well.

Otherwise, if  $putIfAbsent(k)$  performs a successful CAS of handle index from  $\perp$  (line 38), then  $putIfAbsent(k)$  returns true and by Definition 1 there was no handle associated with  $k$  just before the CAS. By Corollary 5, in  $lin(r)$  the map does not hold  $k$ , and since this is the l.p. of  $putIfAbsent$  then the corresponding  $putIfAbsent(k)$  in  $lin(r)$  returns true as required.  $\square$

**Claim 8 (ComputeIfPresent).** In run  $r$ , if  $computeIfPresent(k)$  returns true then in  $lin(r)$  the corresponding  $computeIfPresent(k)$  returns true, and if  $computeIfPresent(k)$  returns false then the corresponding  $computeIfPresent(k)$  in  $lin(r)$  returns false.

*Proof.* If  $computeIfPresent(k)$  finds a non-deleted handle  $h$  and there is a successful nested call to handle compute (line 50), then  $computeIfPresent(k)$  returns true and by Claim 2  $h$  is associated with  $k$ . By Corollary 5, in  $lin(r)$   $k$  is inserted and not removed (the map holds  $k$ ) and since this is the l.p. of  $computeIfPresent$  then the corresponding  $computeIfPresent(k)$  in  $lin(r)$  returns true as well.

If  $lookup(k)$  by  $computeIfPresent(k)$  returns  $\perp$ , or if  $computeIfPresent(k)$  reads that the handle index is  $\perp$  (line 47), then there is no handle associated with key  $k$ , and  $computeIfPresent(k)$  returns false. By Corollary 5, in  $lin(r)$  the map does not hold  $k$ , and since this is the l.p. of  $computeIfPresent$  then the corresponding  $computeIfPresent(k)$  in  $lin(r)$  returns false as required.

Otherwise, a successful CAS of handle index to  $\perp$  is performed by  $computeIfPresent(k)$  (line 56), from a handle index pointing to a deleted handle (line 49). Then  $computeIfPresent(k)$  returns false and by Definition 1 there is no handle associated with  $k$  just before the CAS and right after it. By Corollary 5, in  $lin(r)$  the map does not hold  $k$ , and since this is the l.p. of  $computeIfPresent$  then the corresponding  $computeIfPresent(k)$  in  $lin(r)$  returns false.  $\square$

**Claim 9 (Put).** In run  $r$ , if  $put(k)$  inserts  $k$  and returns then in  $lin(r)$  the corresponding  $put(k)$  inserts  $k$  and returns, and if  $put(k)$  replaces  $k$ 's value and returns then in  $lin(r)$  the corresponding  $put(k)$  replaces  $k$ 's value and returns.

*Proof.* If  $put(k)$  finds a non-deleted handle  $h$  and there is a successful nested call to handle put (line 23), then  $put(k)$  replaces  $k$ 's value and returns, and by Claim 2  $h$  is associated with  $k$ . By Corollary 5, in  $lin(r)$  the map holds  $k$  and since this is the l.p. of  $put$  then the corresponding  $put(k)$  in  $lin(r)$  replaces  $k$ 's value and returns as well.

Otherwise,  $\text{put}(k)$  performs a successful CAS of handle index (line 38) from  $\perp$ , and inserts  $k$  and returns. By Definition 1 there is no handle associated with  $k$  just before the CAS, and there is one right after the CAS (the handle is initialized as non-deleted). Since this is the l.p. of put, and by Corollary 5 in  $\text{lin}(r)$  the map does not hold  $k$  before the l.p. and does after. Therefore, the corresponding  $\text{put}(k)$  in  $\text{lin}(r)$  inserts  $k$  and returns as required.  $\square$

**Claim 10** ( $\text{PutIfAbsentComputeIfPresent}$ ). In run  $r$ , if  $\text{putIfAbsentComputeIfPresent}(k)$  inserts  $k$  and returns then in  $\text{lin}(r)$  the corresponding  $\text{putIfAbsentComputeIfPresent}(k)$  inserts  $k$  and returns, and if  $\text{putIfAbsentComputeIfPresent}(k)$  updates  $k$ 's value and returns then in  $\text{lin}(r)$  the corresponding  $\text{putIfAbsentComputeIfPresent}(k)$  updates  $k$ 's value and returns.

*Proof.* If  $\text{putIfAbsentComputeIfPresent}(k)$  performs a successful CAS of handle index (line 38) from  $\perp$ , then it inserts  $k$  and returns. By Definition 1 there is no handle associated with  $k$  just before the CAS, and there is one right after the CAS (the handle is initialized as non-deleted). Since this is the l.p. of  $\text{putIfAbsentComputeIfPresent}$ , and by Corollary 5 in  $\text{lin}(r)$  the map does not hold  $k$  before the l.p. and does after. Therefore, the corresponding  $\text{putIfAbsentComputeIfPresent}(k)$  in  $\text{lin}(r)$  inserts  $k$  and returns as required.

Otherwise,  $\text{putIfAbsentComputeIfPresent}(k)$  finds a non-deleted handle  $h$  and there is a successful nested call to handle compute (line 25), then  $\text{putIfAbsentComputeIfPresent}(k)$  updates  $k$ 's value and returns, and by Claim 2  $h$  is associated with  $k$ . By Corollary 5, in  $\text{lin}(r)$  the map holds  $k$  and since this is the l.p. of  $\text{putIfAbsentComputeIfPresent}$  then the corresponding  $\text{putIfAbsentComputeIfPresent}(k)$  in  $\text{lin}(r)$  updates  $k$ 's value and returns as well.  $\square$

**Claim 11** ( $\text{Remove}$ ). In run  $r$ , if  $\text{remove}(k)$  removes  $k$  and returns then in  $\text{lin}(r)$  the corresponding  $\text{remove}(k)$  removes  $k$  and returns, and if  $\text{remove}(k)$  returns unsuccessfully (without removing any key) then in  $\text{lin}(r)$  the corresponding  $\text{remove}(k)$  returns unsuccessfully.

*Proof.* If  $\text{remove}(k)$  finds a non-deleted handle  $h$  and a successful nested call to handle remove occurs, setting the handle to deleted (line 52), then  $\text{remove}(k)$  removes  $k$  and returns. By Claim 2  $h$  is associated with  $k$  before and there is no handle associated with  $k$  right after (by Definition 1). Since this is the l.p. of remove, and by Corollary 5 in  $\text{lin}(r)$  the map does hold  $k$  before the l.p. and does not after. Therefore, the corresponding  $\text{remove}(k)$  in  $\text{lin}(r)$  removes  $k$  and returns as required.

If  $\text{lookUp}(k)$  by  $\text{remove}(k)$  returns  $\perp$ , or if  $\text{remove}(k)$  reads that the handle index is  $\perp$  (line 47), then there is no handle associated with key  $k$ , and  $\text{remove}(k)$  returns unsuccessfully. By Corollary 5, in  $\text{lin}(r)$  the map does not hold  $k$ , and since this is the l.p. of remove then the corresponding  $\text{remove}(k)$  in  $\text{lin}(r)$  returns unsuccessfully as required.

Otherwise, a successful CAS of handle index to  $\perp$  is performed by  $\text{remove}(k)$  (line 56), from a handle index pointing to a deleted handle (line 49). Then  $\text{remove}(k)$  returns and by Definition 1 there is no handle associated with  $k$  just before the CAS and right after it. By Corollary 5, in  $\text{lin}(r)$  the map does not hold  $k$ , and since this is the l.p. of remove then the corresponding  $\text{remove}(k)$  in  $\text{lin}(r)$  returns unsuccessfully.  $\square$

Having shown that all of Oak's operations behave the same way in a run  $r$  and its linearization  $\text{lin}(r)$ , we can conclude the following theorem:

**Theorem 12.** Oak is linearizable with the l.p.s defined in Section 4.5.